

# HW5

Yue Zhang

2025-10-19

```
setwd("/Users/yuezhang/Documents/Biostat/PH1976")
getwd()
library(ISLR2)
library(ggplot2)
library(tidyr)
library(dplyr)
library(tidyverse)

## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
## v dplyr      1.1.4      v readr      2.1.5
## v forcats    1.0.0      v stringr   1.5.1
## v lubridate  1.9.3      v tibble    3.2.1
## v purrr      1.1.0
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors

library(ggcorrplot)
library(MASS)

##
## Attaching package: 'MASS'
##
## The following object is masked from 'package:dplyr':
##
##   select
##
## The following object is masked from 'package:ISLR2':
##
##   Boston

library(class)
library(glmnet)

## Loading required package: Matrix
##
## Attaching package: 'Matrix'
##
## The following objects are masked from 'package:tidyr':
##
##   expand, pack, unpack
##
## Loaded glmnet 4.1-10
```

```
library(leaps)
library(pls)
```

```
##
## Attaching package: 'pls'
##
## The following object is masked from 'package:stats':
##
##      loadings
```

3.

- (a). As  $s$  increases from 0, the training RSS will steadily decrease as the model becomes more flexible.
- (b). For test RSS, it will decrease initially and then eventually starts increasing in a U shape because if we fit too many independent variables, the model might be overfitting.
- (c). The variance will steadily increase as more variables are included, the model becomes complex and sensitive, it will change drastically between different datasets.
- (d). The bias will steadily decrease because we include more variable and more noises, the difference between the average prediction of the model and the true value will be smaller.
- (e). The irreducible error will remain constant because it is caused by inherent randomness or noise in the data and it is unavoidable.

5.

(a).

$$L(\beta_1, \beta_2) = \sum_{i=1}^2 (y_i - (\beta_1 + \beta_2)x_i - \beta_0)^2 + \lambda(\beta_1^2 + \beta_2^2).$$

As  $x_{11} = x_{12}$ ,  $x_{21} = x_{22}$ , we can say that  $x_1 = x_{11} = x_{12}$ ,  $x_2 = x_{21} = x_{22}$ . As  $x_{11} + x_{21} = 0$ ,  $x_{21} + x_{22} = 0$ ,  $x_1 = -x_2$ . Therefore,  $\beta_0 = 0$ , we need to minimize:

$$L(\beta_1, \beta_2) = \sum_{i=1}^2 (y_i - (\beta_1 + \beta_2)x_i)^2 + \lambda(\beta_1^2 + \beta_2^2).$$

(b).

$$\frac{\partial L}{\partial \beta_1} = -2 \sum_{i=1}^2 x_i (y_i - (\beta_1 + \beta_2)x_i) + 2\lambda\beta_1,$$

$$\frac{\partial L}{\partial \beta_2} = -2 \sum_{i=1}^2 x_i (y_i - (\beta_1 + \beta_2)x_i) + 2\lambda\beta_2.$$

To minimize these equations, we need to set them as zero

$$\hat{\beta}_1 = \hat{\beta}_2 = \frac{\sum_{i=1}^2 x_i (y_i - (\beta_1 + \beta_2)x_i)}{\lambda}$$

(c).

$$L(\beta_1, \beta_2) = \sum_{i=1}^2 (y_i - (\beta_1 + \beta_2)x_i)^2 + \lambda(|\beta_1| + |\beta_2|).$$

(d). Assumes that there are some  $\alpha$  value that satisfies  $\beta_1 + \beta_2 = \alpha$ ,  $|\beta_1| + |\beta_2| \geq |\alpha|$  Therefore,

$$L(\beta_1, \beta_2) \geq \sum_{i=1}^2 (y_i - \alpha x_i)^2 + \lambda\alpha.$$

8.

```
set.seed(1215)
X = rnorm(100)
epsilon = rnorm(100)
```

$$Y = 8 + 2 * X + 3 * X^2 + 4 * X^3 + \text{epsilon}$$

```
## Subset selection object
## Call: regsubsets.formula(data = df, Y ~ poly(X, 10, raw = TRUE), nvmax = 10)
## 10 Variables   (and intercept)
##               Forced in Forced out
## poly(X, 10, raw = TRUE)1      FALSE      FALSE
## poly(X, 10, raw = TRUE)2      FALSE      FALSE
## poly(X, 10, raw = TRUE)3      FALSE      FALSE
## poly(X, 10, raw = TRUE)4      FALSE      FALSE
## poly(X, 10, raw = TRUE)5      FALSE      FALSE
## poly(X, 10, raw = TRUE)6      FALSE      FALSE
## poly(X, 10, raw = TRUE)7      FALSE      FALSE
## poly(X, 10, raw = TRUE)8      FALSE      FALSE
## poly(X, 10, raw = TRUE)9      FALSE      FALSE
## poly(X, 10, raw = TRUE)10     FALSE      FALSE
## 1 subsets of each size up to 10
## Selection Algorithm: exhaustive
##           poly(X, 10, raw = TRUE)1 poly(X, 10, raw = TRUE)2
## 1  ( 1 ) " "                      " "
## 2  ( 1 ) " "                      "*"
## 3  ( 1 ) "*"                       "*"
## 4  ( 1 ) "*"                       "*"
## 5  ( 1 ) "*"                       "*"
## 6  ( 1 ) "*"                       "*"
## 7  ( 1 ) "*"                       "*"
## 8  ( 1 ) "*"                       "*"
## 9  ( 1 ) "*"                       "*"
## 10 ( 1 ) "*"                      "*"
##           poly(X, 10, raw = TRUE)3 poly(X, 10, raw = TRUE)4
## 1  ( 1 ) "*"                      " "
## 2  ( 1 ) "*"                      " "
## 3  ( 1 ) "*"                      " "
## 4  ( 1 ) "*"                      " "
## 5  ( 1 ) "*"                      " "
## 6  ( 1 ) "*"                      " "
## 7  ( 1 ) "*"                      "*"
## 8  ( 1 ) "*"                      "*"

```

```
## 9 ( 1 ) "*" "*"
## 10 ( 1 ) "*" "*"
##      poly(X, 10, raw = TRUE)5 poly(X, 10, raw = TRUE)6
## 1 ( 1 ) " " " "
## 2 ( 1 ) " " " "
## 3 ( 1 ) " " " "
## 4 ( 1 ) " " "*"
## 5 ( 1 ) " " " "
## 6 ( 1 ) " " "*"
## 7 ( 1 ) " " "*"
## 8 ( 1 ) " " "*"
## 9 ( 1 ) " " "*"
## 10 ( 1 ) "*" "*"
##      poly(X, 10, raw = TRUE)7 poly(X, 10, raw = TRUE)8
## 1 ( 1 ) " " " "
## 2 ( 1 ) " " " "
## 3 ( 1 ) " " " "
## 4 ( 1 ) " " " "
## 5 ( 1 ) " " "*"
## 6 ( 1 ) " " "*"
## 7 ( 1 ) " " "*"
## 8 ( 1 ) " " "*"
## 9 ( 1 ) "*" "*"
## 10 ( 1 ) "*" "*"
##      poly(X, 10, raw = TRUE)9 poly(X, 10, raw = TRUE)10
## 1 ( 1 ) " " " "
## 2 ( 1 ) " " " "
## 3 ( 1 ) " " " "
## 4 ( 1 ) " " " "
## 5 ( 1 ) " " "*"
## 6 ( 1 ) " " "*"
## 7 ( 1 ) " " "*"
## 8 ( 1 ) "*" "*"
## 9 ( 1 ) "*" "*"
## 10 ( 1 ) "*" "*"

```

```
min_Cp = which.min(summary1$cp)
min_BIC = which.min(summary1$bic)
max_adjR = which.max(summary1$adjr2)

cat(
  "Min Mallow's CP:", min_Cp, "\n",
  "Min BIC:", min_BIC, "\n",
  "Max Adjusted R-squared:", max_adjR, sep = " ")

```

```
## Min Mallow's CP:3
## Min BIC:3
## Max Adjusted R-squared:3

```

```
coef(best_sub, min_Cp)

```

```
##      (Intercept) poly(X, 10, raw = TRUE)1 poly(X, 10, raw = TRUE)2
##      8.077001      1.965296      2.922793
## poly(X, 10, raw = TRUE)3
##      4.038280

```

```
coef(best_sub, min_BIC)
```


```
##              (Intercept) poly(X, 10, raw = TRUE)1 poly(X, 10, raw = TRUE)2  
##              8.077001          1.965296          2.922793  
## poly(X, 10, raw = TRUE)3  
##              4.038280
```

```
coef(best_sub, max_adjR)
```


```
##              (Intercept) poly(X, 10, raw = TRUE)1 poly(X, 10, raw = TRUE)2  
##              8.077001          1.965296          2.922793  
## poly(X, 10, raw = TRUE)3  
##              4.038280
```

```
par(mfrow = c(1,1))
```


```
plot(1:10, summary1$cp, type = "b", xlab = "Model size", ylab = "Mallow's Cp")  
abline(v = min_Cp, lty = 2)
```



```
plot(1:10, summary1$bic, type = "b", xlab = "Model size", ylab = "BIC")  
abline(v = min_BIC, lty = 2)
```



```
plot(1:10, summary1$adjr2, type = "b", xlab = "Model size", ylab = "Adjusted R-squared")
abline(v = max_adjR, lty = 2)
```



on the results, the best model is:  $Y = 8.0770 + 1.9653X + 2.9228X^2 + 4.0383X^3$ .

Based

(d). #Forward

```
forward_sub = regsubsets(data = df, Y ~ poly(X, 10, raw = TRUE), nvmax = 10, method = "forward")
summary2 = summary(forward_sub)
summary2
```

```
## Subset selection object
```

```
## Call: regsubsets.formula(data = df, Y ~ poly(X, 10, raw = TRUE), nvmax = 10,
```

```

##      method = "forward")
## 10 Variables (and intercept)
##              Forced in Forced out
## poly(X, 10, raw = TRUE)1      FALSE      FALSE
## poly(X, 10, raw = TRUE)2      FALSE      FALSE
## poly(X, 10, raw = TRUE)3      FALSE      FALSE
## poly(X, 10, raw = TRUE)4      FALSE      FALSE
## poly(X, 10, raw = TRUE)5      FALSE      FALSE
## poly(X, 10, raw = TRUE)6      FALSE      FALSE
## poly(X, 10, raw = TRUE)7      FALSE      FALSE
## poly(X, 10, raw = TRUE)8      FALSE      FALSE
## poly(X, 10, raw = TRUE)9      FALSE      FALSE
## poly(X, 10, raw = TRUE)10     FALSE      FALSE
## 1 subsets of each size up to 10
## Selection Algorithm: forward
##      poly(X, 10, raw = TRUE)1 poly(X, 10, raw = TRUE)2
## 1 ( 1 ) " " " "
## 2 ( 1 ) " " "*"
## 3 ( 1 ) "*" "*"
## 4 ( 1 ) "*" "*"
## 5 ( 1 ) "*" "*"
## 6 ( 1 ) "*" "*"
## 7 ( 1 ) "*" "*"
## 8 ( 1 ) "*" "*"
## 9 ( 1 ) "*" "*"
## 10 ( 1 ) "*" "*"
##      poly(X, 10, raw = TRUE)3 poly(X, 10, raw = TRUE)4
## 1 ( 1 ) "*" " "
## 2 ( 1 ) "*" " "
## 3 ( 1 ) "*" " "
## 4 ( 1 ) "*" " "
## 5 ( 1 ) "*" "*"
## 6 ( 1 ) "*" "*"
## 7 ( 1 ) "*" "*"
## 8 ( 1 ) "*" "*"
## 9 ( 1 ) "*" "*"
## 10 ( 1 ) "*" "*"
##      poly(X, 10, raw = TRUE)5 poly(X, 10, raw = TRUE)6
## 1 ( 1 ) " " " "
## 2 ( 1 ) " " " "
## 3 ( 1 ) " " " "
## 4 ( 1 ) " " "*"
## 5 ( 1 ) " " "*"
## 6 ( 1 ) " " "*"
## 7 ( 1 ) " " "*"
## 8 ( 1 ) " " "*"
## 9 ( 1 ) " " "*"
## 10 ( 1 ) "*" "*"
##      poly(X, 10, raw = TRUE)7 poly(X, 10, raw = TRUE)8
## 1 ( 1 ) " " " "
## 2 ( 1 ) " " " "
## 3 ( 1 ) " " " "
## 4 ( 1 ) " " " "
## 5 ( 1 ) " " " "

```

```

## 6 ( 1 ) " " " "
## 7 ( 1 ) " " "*"
## 8 ( 1 ) " " "*"
## 9 ( 1 ) "*" "*"
## 10 ( 1 ) "*" "*"
##          poly(X, 10, raw = TRUE)9 poly(X, 10, raw = TRUE)10
## 1 ( 1 ) " " " "
## 2 ( 1 ) " " " "
## 3 ( 1 ) " " " "
## 4 ( 1 ) " " " "
## 5 ( 1 ) " " " "
## 6 ( 1 ) " " "*"
## 7 ( 1 ) " " "*"
## 8 ( 1 ) "*" "*"
## 9 ( 1 ) "*" "*"
## 10 ( 1 ) "*" "*"

min_Cp2 = which.min(summary2$cp)
min_BIC2 = which.min(summary2$bic)
max_adjR2 = which.max(summary2$adjr2)

cat(
  "Min Mallow's CP:", min_Cp2, "\n",
  "Min BIC:", min_BIC2, "\n",
  "Max Adjusted R-squared:", max_adjR2, sep = "")

## Min Mallow's CP:3
## Min BIC:3
## Max Adjusted R-squared:3

coef(forward_sub, min_Cp2)

##          (Intercept) poly(X, 10, raw = TRUE)1 poly(X, 10, raw = TRUE)2
##          8.077001          1.965296          2.922793
## poly(X, 10, raw = TRUE)3
##          4.038280

coef(forward_sub, min_BIC2)

##          (Intercept) poly(X, 10, raw = TRUE)1 poly(X, 10, raw = TRUE)2
##          8.077001          1.965296          2.922793
## poly(X, 10, raw = TRUE)3
##          4.038280


coef(forward_sub, max_adjR2)

##          (Intercept) poly(X, 10, raw = TRUE)1 poly(X, 10, raw = TRUE)2
##          8.077001          1.965296          2.922793
## poly(X, 10, raw = TRUE)3
##          4.038280


par(mfrow = c(1,1))
plot(1:10, summary2$cp, type = "b", xlab = "Model size", ylab = "Mallow's Cp")
abline(v = min_Cp2, lty = 2)

```






```
plot(1:10, summary2$bic, type = "b", xlab = "Model size", ylab = "BIC")
abline(v = min_BIC2, lty = 2)
```



```
plot(1:10, summary2$adjr2, type = "b", xlab = "Model size", ylab = "Adjusted R-squared")
abline(v = max_adjR2, lty = 2)
```



Model size

The

best forward subset model is:  $Y = 8.0770 + 1.9653X + 2.9228X^2 + 4.0383X^3$  which is the same as part c.

#Backward

```
backward_sub = regsubsets(data = df, Y ~ poly(X, 10, raw = TRUE), nvmax = 10, method = "backward")
summary3 = summary(backward_sub)
summary3
```

```
## Subset selection object
## Call: regsubsets.formula(data = df, Y ~ poly(X, 10, raw = TRUE), nvmax = 10,
##      method = "backward")
## 10 Variables (and intercept)
```

	Forced in	Forced out
## poly(X, 10, raw = TRUE)1	FALSE	FALSE
## poly(X, 10, raw = TRUE)2	FALSE	FALSE
## poly(X, 10, raw = TRUE)3	FALSE	FALSE
## poly(X, 10, raw = TRUE)4	FALSE	FALSE
## poly(X, 10, raw = TRUE)5	FALSE	FALSE
## poly(X, 10, raw = TRUE)6	FALSE	FALSE
## poly(X, 10, raw = TRUE)7	FALSE	FALSE
## poly(X, 10, raw = TRUE)8	FALSE	FALSE
## poly(X, 10, raw = TRUE)9	FALSE	FALSE
## poly(X, 10, raw = TRUE)10	FALSE	FALSE

## 1 subsets of each size up to 10

## Selection Algorithm: backward

	poly(X, 10, raw = TRUE)1	poly(X, 10, raw = TRUE)2
## 1 ( 1 )	" "	" "
## 2 ( 1 )	" "	"*"
## 3 ( 1 )	"*"	"*"
## 4 ( 1 )	"*"	"*"
## 5 ( 1 )	"*"	"*"
## 6 ( 1 )	"*"	"*"
## 7 ( 1 )	"*"	"*"
## 8 ( 1 )	"*"	"*"

```

## 9 ( 1 ) "*" "*"
## 10 ( 1 ) "*" "*"
##      poly(X, 10, raw = TRUE)3 poly(X, 10, raw = TRUE)4
## 1 ( 1 ) "*" " "
## 2 ( 1 ) "*" " "
## 3 ( 1 ) "*" " "
## 4 ( 1 ) "*" " "
## 5 ( 1 ) "*" " "
## 6 ( 1 ) "*" " "
## 7 ( 1 ) "*" "*"
## 8 ( 1 ) "*" "*"
## 9 ( 1 ) "*" "*"
## 10 ( 1 ) "*" "*"
##      poly(X, 10, raw = TRUE)5 poly(X, 10, raw = TRUE)6
## 1 ( 1 ) " " " "
## 2 ( 1 ) " " " "
## 3 ( 1 ) " " " "
## 4 ( 1 ) " " " "
## 5 ( 1 ) " " " "
## 6 ( 1 ) " " "*"
## 7 ( 1 ) " " "*"
## 8 ( 1 ) " " "*"
## 9 ( 1 ) " " "*"
## 10 ( 1 ) "*" "*"
##      poly(X, 10, raw = TRUE)7 poly(X, 10, raw = TRUE)8
## 1 ( 1 ) " " " "
## 2 ( 1 ) " " " "
## 3 ( 1 ) " " " "
## 4 ( 1 ) " " "*"
## 5 ( 1 ) " " "*"
## 6 ( 1 ) " " "*"
## 7 ( 1 ) " " "*"
## 8 ( 1 ) " " "*"
## 9 ( 1 ) "*" "*"
## 10 ( 1 ) "*" "*"
##      poly(X, 10, raw = TRUE)9 poly(X, 10, raw = TRUE)10
## 1 ( 1 ) " " " "
## 2 ( 1 ) " " " "
## 3 ( 1 ) " " " "
## 4 ( 1 ) " " " "
## 5 ( 1 ) " " "*"
## 6 ( 1 ) " " "*"
## 7 ( 1 ) " " "*"
## 8 ( 1 ) "*" "*"
## 9 ( 1 ) "*" "*"
## 10 ( 1 ) "*" "*"

```

```

min_Cp3 = which.min(summary3$cp)
min_BIC3 = which.min(summary3$bic)
max_adjR3 = which.max(summary3$adjr2)

cat(
  "Min Mallow's CP:", min_Cp3, "\n",
  "Min BIC:", min_BIC3, "\n",

```

```

"Max Adjusted R-squared:", max_adjR3, sep = "")


## Min Mallow's CP:3
## Min BIC:3
## Max Adjusted R-squared:3
coef(backward_sub, min_Cp3)

##           (Intercept) poly(X, 10, raw = TRUE)1 poly(X, 10, raw = TRUE)2
##           8.077001          1.965296          2.922793
## poly(X, 10, raw = TRUE)3
##           4.038280
coef(backward_sub, min_BIC3)

##           (Intercept) poly(X, 10, raw = TRUE)1 poly(X, 10, raw = TRUE)2
##           8.077001          1.965296          2.922793
## poly(X, 10, raw = TRUE)3
##           4.038280
coef(backward_sub, max_adjR3)

##           (Intercept) poly(X, 10, raw = TRUE)1 poly(X, 10, raw = TRUE)2
##           8.077001          1.965296          2.922793
## poly(X, 10, raw = TRUE)3
##           4.038280
par(mfrow = c(1,1))
plot(1:10, summary3$cp, type = "b", xlab = "Model size", ylab = "Mallow's Cp")
abline(v = min_Cp3, lty = 2)


```




```

plot(1:10, summary3$bic, type = "b", xlab = "Model size", ylab = "BIC")
abline(v = min_BIC3, lty = 2)

```



```
plot(1:10, summary3$adjr2, type = "b", xlab = "Model size", ylab = "Adjusted R-squared")
abline(v = max_adjR3, lty = 2)
```




The best backward subset model is:  $Y = 8.0770 + 1.9653X + 2.9228X^2 + 4.0383X^3$  which is the same as part c. (e).

```
lasso = cv.glmnet(poly(df$X, 10, raw = TRUE), df$Y, alpha = 1, nfolds = 10, standardize = TRUE)
best_lambda = lasso$lambda.min
best_lambda
```

```
## [1] 0.1437779
```

```
plot(lasso)
```



```
coef_min = coef(lasso, s = "lambda.min")
coef_min
```

```
## 11 x 1 sparse Matrix of class "dgCMatrix"
##           lambda.min
## (Intercept)  8.171832
## 1           1.816136
## 2           2.852118
## 3           4.052946
## 4           .
## 5           .
## 6           .
## 7           .
## 8           .
## 9           .
## 10          .
```

The best backward subset model is:  $Y = 8.1718 + 1.8161X + 2.8521X^2 + 4.0529X^3$  which is slightly different as part c.

(f). #Best Subset

```
Y2 = 8 + 7 * X^7 + epsilon
df2 = data.frame(X, Y2)
```

```
best_sub2 = regsubsets(data = df2, Y2 ~ poly(X, 10, raw = TRUE), nvmax = 10)
summary4 = summary(best_sub2)
summary4
```

```
## Subset selection object
```

```

## Call: regsubsets.formula(data = df2, Y2 ~ poly(X, 10, raw = TRUE),
##      nvmax = 10)
## 10 Variables (and intercept)
##
##              Forced in Forced out
## poly(X, 10, raw = TRUE)1      FALSE      FALSE
## poly(X, 10, raw = TRUE)2      FALSE      FALSE
## poly(X, 10, raw = TRUE)3      FALSE      FALSE
## poly(X, 10, raw = TRUE)4      FALSE      FALSE
## poly(X, 10, raw = TRUE)5      FALSE      FALSE
## poly(X, 10, raw = TRUE)6      FALSE      FALSE
## poly(X, 10, raw = TRUE)7      FALSE      FALSE
## poly(X, 10, raw = TRUE)8      FALSE      FALSE
## poly(X, 10, raw = TRUE)9      FALSE      FALSE
## poly(X, 10, raw = TRUE)10     FALSE      FALSE
## 1 subsets of each size up to 10
## Selection Algorithm: exhaustive
##      poly(X, 10, raw = TRUE)1 poly(X, 10, raw = TRUE)2
## 1  ( 1 ) " " " "
## 2  ( 1 ) " " " "
## 3  ( 1 ) " " " "
## 4  ( 1 ) " " " "
## 5  ( 1 ) " " " "
## 6  ( 1 ) " " " "
## 7  ( 1 ) "*" " "
## 8  ( 1 ) "*" "*"
## 9  ( 1 ) "*" "*"
## 10 ( 1 ) "*" "*"
##      poly(X, 10, raw = TRUE)3 poly(X, 10, raw = TRUE)4
## 1  ( 1 ) " " " "
## 2  ( 1 ) " " " "
## 3  ( 1 ) " " "*"
## 4  ( 1 ) " " " "
## 5  ( 1 ) "*" " "
## 6  ( 1 ) " " "*"
## 7  ( 1 ) " " "*"
## 8  ( 1 ) " " "*"
## 9  ( 1 ) "*" "*"
## 10 ( 1 ) "*" "*"
##      poly(X, 10, raw = TRUE)5 poly(X, 10, raw = TRUE)6
## 1  ( 1 ) " " " "
## 2  ( 1 ) " " " "
## 3  ( 1 ) "*" " "
## 4  ( 1 ) "*" " "
## 5  ( 1 ) "*" " "
## 6  ( 1 ) " " "*"
## 7  ( 1 ) "*" "*"
## 8  ( 1 ) "*" "*"
## 9  ( 1 ) " " "*"
## 10 ( 1 ) "*" "*"
##      poly(X, 10, raw = TRUE)7 poly(X, 10, raw = TRUE)8
## 1  ( 1 ) "*" " "
## 2  ( 1 ) "*" "*"
## 3  ( 1 ) "*" " "
## 4  ( 1 ) "*" "*"

```

```

## 5 ( 1 ) "*" "*"
## 6 ( 1 ) "*" "*"
## 7 ( 1 ) "*" "*"
## 8 ( 1 ) "*" "*"
## 9 ( 1 ) "*" "*"
## 10 ( 1 ) "*" "*"
##          poly(X, 10, raw = TRUE)9 poly(X, 10, raw = TRUE)10
## 1 ( 1 ) " " " "
## 2 ( 1 ) " " " "
## 3 ( 1 ) " " " "
## 4 ( 1 ) "*" " "
## 5 ( 1 ) "*" " "
## 6 ( 1 ) "*" "*"
## 7 ( 1 ) " " "*"
## 8 ( 1 ) " " "*"
## 9 ( 1 ) "*" "*"
## 10 ( 1 ) "*" "*"

min_Cp4 = which.min(summary4$cp)
min_BIC4 = which.min(summary4$bic)
max_adjR4 = which.max(summary4$adjr2)

cat(
  "Min Mallow's CP:", min_Cp4, "\n",
  "Min BIC:", min_BIC4, "\n",
  "Max Adjusted R-squared:", max_adjR4, sep = "")

## Min Mallow's CP:1
## Min BIC:1
## Max Adjusted R-squared:2

coef(best_sub2, min_Cp4)

##          (Intercept) poly(X, 10, raw = TRUE)7
##          8.009195          7.000194

coef(best_sub2, min_BIC4)

##          (Intercept) poly(X, 10, raw = TRUE)7
##          8.009195          7.000194


coef(best_sub2, max_adjR4)

##          (Intercept) poly(X, 10, raw = TRUE)7 poly(X, 10, raw = TRUE)8
##          8.0315273441          7.0012992425          -0.0004521743


par(mfrow = c(1,1))
plot(1:10, summary4$cp, type = "b", xlab = "Model size", ylab = "Mallow's Cp")
abline(v = min_Cp4, lty = 2)

```






```
plot(1:10, summary4$bic, type = "b", xlab = "Model size", ylab = "BIC")
abline(v = min_BIC4, lty = 2)
```



```
plot(1:10, summary4$adjr2, type = "b", xlab = "Model size", ylab = "Adjusted R-squared")
abline(v = max_adjR4, lty = 2)
```




Based on the Mallows's Cp and BIC, the best subset model should be:  $Y = 8.0092 + 7.0001X^7$ , if based on the adjusted R-squared, the best subset model is:  $Y = 8.0315 + 7.0013X^7 - 0.0005X^8$

#Lasso

```
lasso2 = cv.glmnet(poly(df2$X, 10, raw = TRUE), df2$Y2, alpha = 1, nfolds = 10, standardize = TRUE)
best_lambda2 = lasso2$lambda.min
best_lambda2
```

```
## [1] 50.68498
```

```
plot(lasso2)
```



```
coef_min2 = coef(lasso2, s = "lambda.min")
coef_min2
```

```
## 11 x 1 sparse Matrix of class "dgCMatrix"
##          lambda.min
## (Intercept) 19.2461374
## 1          .
## 2          .
## 3          .
## 4          .
## 5          0.2323746
## 6          .
## 7          6.7638456
## 8          .
## 9          .
## 10         .
```

The best Lasso model is:  $Y = 19.2461 + 0.2324X^5 + 6.7638X^7$ .

11.

(a).

```
boston = ISLR2::Boston
head(boston)
```

```
##      crim zn indus chas   nox   rm  age   dis rad tax ptratio lstat medv
## 1 0.00632 18  2.31    0 0.538 6.575 65.2 4.0900  1 296   15.3  4.98 24.0
## 2 0.02731  0  7.07    0 0.469 6.421 78.9 4.9671  2 242   17.8  9.14 21.6
## 3 0.02729  0  7.07    0 0.469 7.185 61.1 4.9671  2 242   17.8  4.03 34.7
## 4 0.03237  0  2.18    0 0.458 6.998 45.8 6.0622  3 222   18.7  2.94 33.4
## 5 0.06905  0  2.18    0 0.458 7.147 54.2 6.0622  3 222   18.7  5.33 36.2
## 6 0.02985  0  2.18    0 0.458 6.430 58.7 6.0622  3 222   18.7  5.21 28.7
```

```

#We will use 10-fold CV RMSE to evaluate the models
y = boston$crim
X_mm = model.matrix(data = boston, crim ~ .)[-1]
p = ncol(X_mm)

pred_regsubsets = function(object, newdata, id){
  form = as.formula(paste("crim ~", paste(names(coef(object, id))[-1], collapse = "+")))
  mm = model.matrix(form, newdata)
  drop(mm %*% coef(object, id))
}

rmse = function(y, yhat) sqrt(mean((y - yhat)^2))

K = 10
fold = sample(rep(1:K, length.out = nrow(boston)))

#Best subsets

best_sub3 = regsubsets(data = boston, crim ~ ., nvmax = p)
summary5 = summary(best_sub3)
summary5

## Subset selection object
## Call: regsubsets.formula(data = boston, crim ~ ., nvmax = p)
## 12 Variables (and intercept)
##           Forced in Forced out
## zn          FALSE      FALSE
## indus       FALSE      FALSE
## chas        FALSE      FALSE
## nox         FALSE      FALSE
## rm          FALSE      FALSE
## age         FALSE      FALSE
## dis         FALSE      FALSE
## rad         FALSE      FALSE
## tax         FALSE      FALSE
## ptratio     FALSE      FALSE
## lstat       FALSE      FALSE
## medv        FALSE      FALSE
## 1 subsets of each size up to 12
## Selection Algorithm: exhaustive
##           zn indus chas nox rm age dis rad tax ptratio lstat medv
## 1 ( 1 ) " " " " " " " " " " " " " " " " " " " " "
## 2 ( 1 ) " " " " " " " " " " " " " " " " " " " "
## 3 ( 1 ) " " " " " " " " " " " " " " " " " " " "
## 4 ( 1 ) "*" " " " " " " " " " " "*" " " " " " " "
## 5 ( 1 ) "*" "*" " " " " " " " " " "*" " " " " " "
## 6 ( 1 ) "*" " " " " " " " " " "*" " " " " "*" "
## 7 ( 1 ) "*" " " " " " " " " " "*" " " " " "*" "
## 8 ( 1 ) "*" "*" " " " " " " " "*" " " " " "*" "
## 9 ( 1 ) "*" "*" " " " " " " " "*" " " " " "*" "
## 10 ( 1 ) "*" " " " " " " " " " "*" " " "*" " "*"
## 11 ( 1 ) "*" "*" " " " " " " " "*" " " "*" " "*"
## 12 ( 1 ) "*" "*" " " " " " " " "*" " " "*" " "*"

```

```
min_Cp5 = which.min(summary5$cp)
min_BIC5 = which.min(summary5$bic)
max_adjR5 = which.max(summary5$adjr2)
```

```
cat(
  "Min Mallow's CP:", min_Cp5, "\n",
  "Min BIC:", min_BIC5, "\n",
  "Max Adjusted R-squared:", max_adjR5, sep = "")
```

```
## Min Mallow's CP:7
## Min BIC:2
## Max Adjusted R-squared:9
```

```
coef(best_sub3, min_Cp5)
```

```
## (Intercept)          zn          nox          dis          rad          ptratio
## 17.4668235    0.0449679 -12.4578238  -0.9425497    0.5615224  -0.3470306
##          lstat          medv
##   0.1147895  -0.1902559
```


```
coef(best_sub3, min_BIC5)
```

```
## (Intercept)          rad          lstat
## -4.3814053    0.5228128    0.2372846
```


```
coef(best_sub3, max_adjR5)
```

```
## (Intercept)          zn          indus          nox          rm          dis
## 13.18247199    0.04305936  -0.08820604 -10.46823468    0.63510592  -1.00637216
##          rad          ptratio          lstat          medv
##   0.56096588  -0.30423849    0.14040855  -0.22012525
```


```
par(mfrow = c(1,1))
plot(1:p, summary5$cp, type = "b", xlab = "Model size", ylab = "Mallow's Cp")
abline(v = min_Cp5, lty = 2)
```



```
plot(1:p, summary5$bic, type = "b", xlab = "Model size", ylab = "BIC")
abline(v = min_BIC5, lty = 2)
```



```
plot(1:p, summary5$adjr2, type = "b", xlab = "Model size", ylab = "Adjusted R-squared")
abline(v = max_adjR5, lty = 2)
```




Based on the Mallows's Cp, the best subset model should be:  $crim = 17.4668 + 0.04497zn - 12.4578nox - 0.9425dis + 0.5615rad - 0.3470ptratio + 0.1148lstat - 0.1903medv$ . Based on the BIC,  $crim = -4.3814 + 0.5228rad + 0.2373lstat$ . For the adjusted R-squared, the best subset model is:  $crim = 13.1825 + 0.0431zn - 0.0882indus - 10.4682nox + 0.6351rm - 1.0064dis + 0.5610rad -$

$0.3042ptratio + 0.1404lstat - 0.2201medv$

#Lasso

```
cv_lasso = cv.glmnet(X_mm, y, alpha = 1, nfolds = 10)
plot(cv_lasso)
```




```
coef(cv_lasso, s = "lambda.min")
```

```
## 13 x 1 sparse Matrix of class "dgCMatrix"
##          lambda.min
## (Intercept) 10.9830172958
## zn          0.0400473944
## indus       -0.0662084982
## chas        -0.7160278330
## nox         -8.0426443608
## rm          0.4946556538
## age         .
## dis        -0.8790279219
## rad         0.5592234693
## tax        -0.0008828103
## ptratio    -0.2541193818
## lstat       0.1364191296
## medv       -0.1937324075
```

Based on the best  $\lambda$ , the best subset model should be:  $crim = 10.9830 + 0.0400zn - 0.0662indus - 0.7160chas - 8.0426nox + 0.4947rm - 0.8790dis + 0.5592rad - 0.0009tax - 0.2541ptratio + 0.1364lstat - 0.1937medv$ .

#Ridge

```
cv_ridge = cv.glmnet(X_mm, y, alpha = 0, nfolds = 10)
plot(cv_ridge)
```



```
coef(cv_ride, s = "lambda.min")
```

```
## 13 x 1 sparse Matrix of class "dgCMatrix"
##          lambda.min
## (Intercept) 5.222257e+00
## zn          3.375080e-02
## indus       -7.773854e-02
## chas        -8.254911e-01
## nox         -4.905342e+00
## rm          5.165813e-01
## age         9.537272e-05
## dis        -7.217088e-01
## rad         4.433556e-01
## tax         3.737247e-03
## ptratio    -1.642366e-01
## lstat       1.563538e-01
## medv       -1.584061e-01
```

The best ridge model is:  $crim = 5.2223 + 0.0338zn - 0.0777indus - 0.8255chas - 4.9053nox + 0.5166rm + 0.0000954age - 0.7217dis + 0.4434rad + 0.00374tax - 0.1642ptratio + 0.1564lstat - 0.1584medv$ .

(b).

```
#Best subsets
cv_bs = matrix(NA_real_, K, p)
for(k in 1:K){
  fit = regsubsets(crim ~ ., data = boston[fold != k, ], nvmax = p)
  for(m in 1:p){
    yhat = pred_regsubsets(fit, boston[fold == k, ], id = m)
    cv_bs[k, m] = sqrt(mean( (boston$crim[fold == k] - yhat)^2 ))
  }
}
```



```

rmse_bs = colMeans(cv_bs)
size_bs = which.min(rmse_bs)
rmse_bs_min = min(rmse_bs)
coef_bs = coef(regsubsets(crim ~ ., data = boston, nvmax = p), size_bs)
coef_bs

##      (Intercept)          zn          indus          chas          nox
## 13.801555544    0.045818028 -0.058345869 -0.828283847 -10.022404078
##           rm          dis          rad          tax          ptratio
##  0.623650191 -1.008539667   0.612822152 -0.003782956 -0.304784434
##          lstat          medv
##  0.137698956 -0.220092318

#Lasso
rmse_lasso = numeric(K)
lambda_lass = numeric(K)

for(k in 1:K){
  tr = fold != k; te = !tr
  fit_cv = cv.glmnet(X_mm[tr, ], y[tr], alpha = 1)
  lambda_lass[k] = fit_cv$lambda.min
  pred = as.numeric(predict(fit_cv, newx = X_mm[te, ], s = "lambda.min"))
  rmse_lasso[k] = rmse(y[te], pred)
}

rmse_lasso_mean = mean(rmse_lasso)

lasso_full = cv.glmnet(X_mm, y, alpha = 1)
coef_lasso = coef(lasso_full, s = "lambda.min")
coef_lasso

```

```

## 13 x 1 sparse Matrix of class "dgCMatrix"
##           lambda.min
## (Intercept)  9.76734200
## zn          0.03781780
## indus       -0.06715706
## chas        -0.67616141
## nox         -7.08411010
## rm          0.43340686
## age         .
## dis        -0.81982181
## rad         0.54122756
## tax         .
## ptratio    -0.23042121
## lstat       0.13565977
## medv       -0.18214662

```

```

#Ridge
rmse_ridge = numeric(K)
lambda_rid = numeric(K)

for(k in 1:K){
  tr = fold != k; te <- !tr
  fit_cv = cv.glmnet(X_mm[tr, ], y[tr], alpha = 0)
  lambda_rid[k] = fit_cv$lambda.min
}

```

```

    pred = as.numeric(predict(fit_cv, newx = X_mm[te, ], s = "lambda.min"))
    rmse_ridge[k] = rmse(y[te], pred)
}

rmse_ridge_mean = mean(rmse_ridge)
ridge_full = cv.glmnet(X_mm, y, alpha = 0)
coef_ridge = coef(ridge_full, s = "lambda.min")
coef_ridge

```

```

## 13 x 1 sparse Matrix of class "dgCMatrix"
##               lambda.min
## (Intercept)  5.22257e+00
## zn           3.37508e-02
## indus        -7.77385e-02
## chas         -8.25491e-01
## nox          -4.90534e+00
## rm           5.16581e-01
## age          9.53727e-05
## dis          -7.21708e-01
## rad          4.43355e-01
## tax          3.73724e-03
## ptratio      -1.64236e-01
## lstat        1.56353e-01
## medv         -1.58406e-01

```

```

c(
  RMSE_best_subset = rmse_bs_min,
  RMSE_lasso       = rmse_lasso_mean,
  RMSE_ridge       = rmse_ridge_mean
)

```

```

## RMSE_best_subset    RMSE_lasso    RMSE_ridge
##           5.707388           5.690106           5.669860

```

Based on the result, we will choose Ridge as the best model as it has the least RMSE. The best model is:  
 $crim = 5.2223 + 0.0338zn - 0.0777indus - 0.8255chas - 4.9053nox + 0.5166rm + 0.0000954age - 0.7217dis + 0.4434rad + 0.00374tax - 0.1642ptratio + 0.1564lstat - 0.1584medv$ .

(c). The ridge model features all the variables. Ridge will shrink coefficients toward 0 but does not set them exactly to 0.